

# MODEL AGGREGATION PRINCIPLES IN FEDERATED LEARNING

Evgeny Osipov



# Introduction to Federated Learning

Federated Learning (FL) is a distributed machine learning paradigm that enables training models across multiple decentralized devices or servers while keeping data localized.

## Key Characteristics:


- Privacy Preservation: Training data remains on local devices
- Decentralized Architecture: No central data repository
- Collaborative Training: Multiple parties contribute to model improvement
- Communication Efficiency: Only model updates are exchanged

# Architecture Overview: Four Core Components

Federated Learning systems consist of four fundamental architectural components.

## Core Components:

1. **Server:** Central coordinator and aggregator
2. **Parties:** Distributed clients/participants
3. **Communication Framework:** Secure data exchange protocols
4. **Aggregation Algorithm:** Model update combination logic



# Architecture Overview: Server

The server orchestrates the FL process, acting as the primary point of control and aggregation.

## Primary Functions:

- Client selection and round management
- Parameter distribution to participating clients
- Model aggregation from client updates
- Performance monitoring and compliance

### SERVER

Central Coordinator & Aggregator

- orchestrates rounds
- distributes global model
- collects client updates

# Architecture Overview: Parties

Parties are the distributed entities/clients that own training data and perform local model training.

## Two Main Categories:

- **Cross-Device FL:** Mobile devices, IoT sensors
- **Cross-Silo FL:** Organizations, companies, institutions

### PARTY 1

- Client/device
- Local data
- Local training

### PARTY 2

- Client/device
- Local data
- Local training

### PARTY N

- Client/device
- Local data
- Local training

# Architecture Overview: Communication framework

The communication framework encompasses protocols for secure, efficient data exchange.

## Core Layers:

- **Transport:** TCP/IP, UDP, WebSockets
- **Security:** TLS/SSL, certificates
- **Application:** gRPC, REST APIs, MQTT

COMMUNICATION FRAMEWORK (secure exchange)  
- authentication / encryption /  
secure channels  
- message formats +  
protocols (client ↔ server)

# FL Categorization Overview

Federated Learning can be categorized based on data distribution patterns across participants:

## 1. Horizontal FL (HFL):

- Same feature spaces, different samples
- Most common form (FedAvg, cross-device)
- Scales to millions of participants

## 2. Vertical FL (VFL):

- Different feature spaces, overlapping samples
- Feature-partitioned data across organizations
- Requires cryptographic techniques

## 3. Federated Transfer Learning:

- Combines FL with transfer learning
- Domain adaptation across heterogeneous environments
- Addresses cold start and data scarcity

# Horizontal Federated Learning (HFL)

Horizontal Federated Learning occurs when participants have identical data schemas but different samples.

## Key Characteristics:

- Feature Space: Same attributes/columns across all participants
- Sample Space: Different/non-overlapping data samples
- Scalability: Naturally scales to large numbers of participants
- Privacy: Strong guarantees with minimal data overlap

## FedAvg Algorithm:

$$w^{t+1} = \sum_{k=1}^K \frac{n_k}{n} w_k^{t+1}$$

# HFL Use Cases

HFL powers many real-world applications.

## Cross-Device FL:

- Gboard: Next-word prediction across Android devices
  - <https://support.google.com/gboard/answer/12373137?hl=en>
- Mobile Apps: Personalized recommendations
- IoT Networks: Smart device coordination

## Cross-Silo FL:

- Banks: Fraud detection collaboration
- Hospitals: Medical research without data sharing
- Companies: Enterprise model training

# HFL Challenges

HFL faces challenges from participant and data diversity.

## Main Challenges:

- Statistical heterogeneity (non-IID data)
- System heterogeneity (varying compute)
- Communication costs and client dropout

# Vertical Federated Learning (VFL)

Vertical Federated Learning occurs when participants have different features for the same entities.

## Key Characteristics:

- Feature Space: Different attributes across participants
- Sample Space: Overlapping users/entities
- Privacy: Very strong (cryptographic techniques required)
- Coordination: Complex sample alignment needed

## Feature Partitioning:

$$x_i = [x_i^{(1)}, x_i^{(2)}, \dots, x_i^{(K)}]$$

Bank A  
 $x_A, y$

$z_A$

Secure Combine  
 $z = z_A + z_B$

$z_B$

Retailer B  
 $x_B$



LULEÅ  
UNIVERSITY  
OF TECHNOLOGY

# VFL Use Cases

VFL enables collaboration across different data silos.

## **Financial Services:**

- Banks share features for credit scoring
- Transaction history + employment data + credit history
- Better risk assessment without data sharing

## **Healthcare:**

- Hospitals combine different medical data
- Lab results + imaging + patient history
- Improved diagnosis with comprehensive data

# VFL Challenges

VFL requires sophisticated cryptographic techniques.

## Main Challenges:

- Sample alignment without sharing identifiers
- Secure multi-party computation
- Homomorphic encryption for joint operations

## Solutions:

- Secret Sharing: Distribute computations
- Homomorphic Encryption: Encrypted arithmetic
- Secure MPC: Privacy-preserving protocols

# Federated Transfer Learning

Type	Users	Features	Example
Horizontal FL	Different	Same	Hospitals in different cities
Vertical FL	Same	Different	Bank + Retail
Federated Transfer Learning	<b>Different</b>	<b>Different</b>	Hospital + Insurance

# Federated Transfer Learning

FL + Transfer Learning Combines federated learning with transfer learning techniques for domain adaptation.

## Key Characteristics:

- Raw data never shared
- Only encrypted embeddings exchanged
- Overlapping ID matching done securely
- Often uses:
  - Secure Multi-Party Computation
  - Homomorphic encryption
  - Differential privacy

# Architectural view

Party A


$x_A \rightarrow f_A - h_A$

Party B

$x_B \rightarrow f_B - h_B$

- $x_A, x_B$  – features (in general non-overlapping)
- $f_A$  – locally trained model to solve organization-specific task
  - E.g. decease forecasting
- $h_A$  – learned feature representation

# Architectural view



After alignment (Even for non-overlapping users):

- Party A (e.g. hospital) model benefits from structural patterns learned by Party B (e.g. Insurer).
- Party B helps shape the embedding space of Party A for better generalization

# Summary so far: FL Categories Comparison

## Use HFL when:

- Large number of participants with similar data structures
- Cross-device applications (mobile, IoT)
- Strong privacy with minimal data overlap


## Use VFL when:

- Few participants with complementary features
- Sensitive data requiring strong privacy guarantees
- Financial services, healthcare applications

## Use Federated Transfer Learning when:

- Participants with different but related domains
- Cold start problems with limited local data
- Evolving environments requiring continual adaptation
- Multi-task learning scenarios

# Recall. What is left?



# What is Model Aggregation - HFL?

Model aggregation is the process of combining individual client models to create a global model in federated learning.

## Goals:

- Leverage collective knowledge from distributed clients
- Maintain data privacy (no raw data exchange)
- Handle client heterogeneity (data, compute, network)
- Ensure robustness against malicious participants
- Optimize for communication efficiency

# Federated Averaging (FedAvg)

## Mathematical Formulation:

$$w^{t+1} = \sum_{k=1}^K \frac{n_k}{n} w_k^{t+1}$$

### Where:

- $w_k^{t+1}$ : Client  $k$ 's model after local training
- $n_k$ : Data samples on client  $k$
- $n = \sum n_k$ : Total samples across clients

# Example FedAvg: Round 1

**Task:** Predict disease risk (binary classification)

**Model:** Single parameter  $w$  (for simplicity)

Initial Model **Server initializes:**  $w^{(0)} = 0.5$

**Step 1:** Server sends  $w^{(0)} = 0.5$  to all clients

**Step 2:** Each client trains locally on their data

**Step 3:** After local training (e.g., 5 epochs of SGD) clients send their trained models back to server

# Example FedAvg: Round 1 (cont.)

**Step 4:** Server computes weighted average

**Formula:**

$$w^{(1)} = \frac{n_A}{n} \cdot w_A^{(1)} + \frac{n_B}{n} \cdot w_B^{(1)} + \frac{n_C}{n} \cdot w_C^{(1)}$$

**Substituting values:**

$$\begin{aligned} w^{(1)} &= \frac{100}{500} \cdot 0.65 + \frac{200}{500} \cdot 0.80 + \frac{200}{500} \cdot 0.75 \\ &= 0.20 \cdot 0.65 + 0.40 \cdot 0.80 + 0.40 \cdot 0.75 \\ &= 0.13 + 0.32 + 0.30 = \mathbf{0.75} \end{aligned}$$

**New global model:**  $w^{(1)} = 0.75$

# Example FedAvg: Round 1 (cont.)

## Compare:

- Hospital A:  $w_A = 0.65$ ) weight: 0.20)
- Hospital B:  $w_B = 0.80$ ) weight: 0.40)
- Hospital C:  $w_C = 0.75$ ) weight: 0.40)
- **Global:**  $w^{(1)} = 0.75$

## Observe:

- Result is closer to B and C (they have more data)
- Hospital A has less influence (only 20% of data)
- **Simple average would give:  $(0.65 + 0.80 + 0.75)/3 = 0.733$**
- Weighted average: 0.75 (accounts for data size)

# Example FedAvg: Round 2 and so on

**Step 1:** Server sends new global model  $w^{(1)} = 0.75$  to all clients

**Step 2:** Each client trains locally starting from  $w^{(1)} = 0.75$

After local training:

**Step 3:** Send updates to server

**Step 4:** Server aggregates:

$$\begin{aligned} w^{(2)} &= 0.20 \cdot 0.78 + 0.40 \cdot 0.82 + 0.40 \cdot 0.80 \\ &= 0.156 + 0.328 + 0.320 = \mathbf{0.804} \end{aligned}$$

**Why weighted is better:**

- Hospitals B and C have  $4 \times$  more data than A
- Their models are trained on more samples (more reliable)
- Weighted average gives them appropriate influence
- Leads to better convergence and final model quality

# FL Security Challenges

Federated Learning introduces distinctive security challenges due to its distributed nature and privacy-preserving design.

**Security threats in FL can undermine both privacy and utility!**

# Attack Impact

Poisoning attacks have cascading effects on FL systems and stakeholders.

## Individual Model Level:

- Incorrect feature engineering
- Regulatory compliance failures
- Debugging based on false explanations

## System-Level Impact:

- Trust erosion in AI systems
- Economic consequences from poor decisions
- Hidden security risks in explanations

## Societal Impact:

- Bias amplification or masking
- Accountability issues
- Flawed policy decisions

# Attack Variants

Attacks can vary based on attacker resources and coordination.

## Single-Client Attacks:

- Limited impact but hard to detect
- Blends with natural data variations

## Multi-Client Coordinated Attacks:

- Sybil attacks with multiple identities
- Combined effect through FedAvg
- Higher impact but easier to detect

## Stealthy Attacks:

- Minimal perturbations
- Adaptive poisoning strategies
- Integration with backdoor attacks

# FL Threat Taxonomy

FL security threats can be categorized into three primary areas.

- 1. Privacy Attacks:** Information leakage from model outputs
- 2. Poisoning Attacks:** Malicious data or model manipulation
- 3. Fairness-Robustness Trade-offs: Balancing multiple objectives**
- 4. Utility-Robustness Trade-offs**

Each category presents unique challenges for FL system design!

# Privacy Risks

Even without sharing raw data, FL systems can leak sensitive information through model outputs and explanations.

## Key Privacy Threats:

- **Membership Inference:** Determine if sample was in training data
- **Property Inference:** Learn statistical properties of training data
- **Explanation Leakage:** Sensitive information through model explanations

# Differential Privacy in FL

Add noise to model updates to provide provable privacy protection.

**Local DP:** Each client adds noise before sending:

$$\tilde{w}_k = w_k + \mathcal{N}(0, \sigma^2)$$

**Central DP:** Server adds noise to aggregated updates:

$$w^{t+1} = \sum_{k=1}^K \frac{n_k}{n} w_k^t + \mathcal{N}(0, \sigma^2)$$

# Poisoning Attacks: Overview

Poisoning attacks involve malicious participants introducing corrupted data or model updates into the FL training process.

## Attack Types:

- **Model Poisoning:** Send manipulated model updates
- **Data Poisoning:** Manipulate local training data
- **Byzantine Attacks:** Arbitrary malicious behavior

Poisoning can target accuracy, explanations, or information extraction!

# Byzantine Threat

A **Byzantine** participant can behave **arbitrarily**:

- send incorrect or inconsistent messages,
- send extreme/malicious gradients/updates,
- collude with others, or remain silent.

In FL, Byzantine clients can:

- break convergence,
- insert backdoors,
- dominate averaging with outliers.

# Model poisoning

A **model poisoning attack** occurs when a malicious client:

- Intentionally modifies its local update
- To manipulate the global model
- While still following the communication protocol

Different from data poisoning (corrupting local data)

# Why FedAvg is Vulnerable

FedAvg assumes:

- Clients are honest.
- Updates reflect legitimate local training.
- Averaging is safe.

But averaging has **no robustness mechanism**.

So a malicious update can:

- Shift the global model
- Inject a backdoor
- Prevent convergence
- Reduce accuracy

# Attacks against FedAvg

With FedAvg:

$$w_{t+1} = w_t + \frac{1}{n} \sum_{i=1}^n \Delta_i .$$

If one client  $m$  is malicious:

$$w_{t+1} = w_t + \frac{1}{n} \left( \sum_{i \neq m} \Delta_i + \Delta_m \right) .$$

A malicious client can scale its update:

$$\Delta_m = n \cdot \Delta^* \Rightarrow w_{t+1} \approx w_t + \Delta^*$$

so a **single client can dominate the round** (“model replacement” style effect).

# Model Poisoning: Common Attack Styles

- **Sign-flip / disruption:**  $\Delta_m = -\alpha \Delta_{\text{honest}}$  to slow or reverse learning.
- **Backdoor:** train locally so that a trigger  $T$  causes prediction  $\hat{y} = y^*$  while keeping normal accuracy high.
- **Model replacement:** craft  $\Delta_m$  to push  $w_{t+1}$  close to an attacker-chosen  $w^*$

# Data Poisoning Against Interpretations

A sophisticated attack that undermines model interpretability by poisoning training data to alter feature importance attributions.

## Core Insight:

Model interpretability methods are sensitive to small, carefully crafted perturbations in training data.

## Key Difference:

- Traditional attacks target model **accuracy**
- This attack targets model **explanations**
- Maintains accuracy while corrupting interpretability

# Attack Methodology

Four-Step Poisoning Process Systematic approach to manipulate model explanations in FL.

**Step 1 - Target Selection:** Choose explanation method (SHAP, LIME, feature importance)

**Step 2 - Poison Generation:** Solve optimization/or craft to create poisoned samples

**Step 3 - Explanation Manipulation:** Feature/Label flip, masking, or inversion attacks

**Step 4 - FL Integration:** Distribute poison across clients, consider FedAvg effects

# Mathematical Formulation

## Objective Function:

$$\min_{\delta} \|\phi(f(\theta, x + \delta)) - \phi_{target}\|^2$$

## Constraints:

$$\|\delta\| \leq \epsilon \quad (\text{perturbation budget})$$

$$\mathcal{L}(f(\theta, x + \delta), y) \leq \tau \quad (\text{accuracy preservation})$$

Trade-off between explanation manipulation and accuracy preservation!

# Manipulation Strategies

Three Attack Variants Different approaches to corrupt model interpretability.

## 1. Feature Flip Attacks:

- Make irrelevant features appear important
- Example: Credit model prioritizes "favorite color" over "credit history"

## 2. Feature Masking Attacks:

- Hide importance of critical features
- Example: Reduce apparent biomarker importance in medical diagnosis

## 3. Explanation Inversion Attacks:

- Make explanations contradict true decision process
- Example: Income-based model appears to prioritize education

# Defense Strategies

FL defense strategies address different attack types with appropriate countermeasures.

## Defense Categories:

- 1. Robust Aggregation:** Defend against malicious updates
- 2. Client Authentication:** Verify participant legitimacy
- 3. Data Validation:** Detect poisoned training data
- 4. Privacy Protection:** Prevent information leakage
- 5. Fairness Mechanisms:** Ensure equitable performance

# Detection and Mitigation

## Statistical Detection:

- Explanation consistency checks
- Outlier detection in explanation patterns
- **Distribution shift monitoring**

## Robust Explanation Methods:

- Ensemble explanations
- Certified explanations with guarantees
- Adversarial training for explanations

# Robust Aggregation Methods: Coordinate-wise Median

Principle: Take median of each parameter coordinate across all client updates.

**Formula:**

$$w_i^{t+1} = \text{median}(\{w_{k,i}^{t+1}\}_{k=1}^K)$$

**Advantages:**

- Robust to value manipulation attacks
- Computationally simple
- No hyperparameters to tune

# Robust Aggregation Methods: Krum Aggregation

Principle Select the update that is closest to other (gradient) updates (multi-dimensional median).

## Algorithm:

1. For each client  $k$ , compute  $\sum_{k' \neq k} \|g_k - g_{k'}\|^2$ , where
2. Select  $k^*$  with minimum distances
3. Average top  $m$  closest updates

$g$ 's are exchanged gradients

# Robust Aggregation Methods: Krum Aggregation

## Setup

- Number of clients:  $n=6$
- Max Byzantine clients:  $f=1$
- Each client sends a **model update vector**  $g$ .
- Krum uses **squared Euclidean distances**.

Krum parameter:

- $m=n-f-2 = 6-1-2 = 3$
- So for each update  $g_i$ , we find its **3 nearest neighbors** (by squared distance) and sum those distances. The update with the **smallest sum** is selected.

# Robust Aggregation Methods: Krum Aggregation

- Honest updates (close to each other):
  - $g_1=[1.00, 1.00]$ ,  $g_2=[1.10, 0.90]$ ,  $g_3=[0.90, 1.20]$ ,  $g_4=[1.05, 1.05]$ ,  $g_5=[0.95, 0.85]$
- Byzantine (malicious outlier):
  - $g_6=[10, -10]$

# Robust Aggregation Methods: Krum Aggregation

Use updates from these nodes

	g1	g2	g3	g4	g5	g6
g1	0	0.020	0.050	0.005	0.025	202.000
g2	0.020	0	0.130	0.025	0.025	198.020
g3	0.050	0.130	0	0.045	0.125	208.250
g4	0.005	0.025	0.045	0	0.050	202.205
g5	0.025	0.025	0.125	0.050	0	199.625
g6	202.000	198.020	208.250	202.205	199.625	0

# Fairness-Robustness Trade-offs

In FL:

- Clients often have **non-IID data**.
- Some clients represent minority populations.
- Their updates may naturally look different from the majority.

Robust aggregation methods (e.g., Krum, trimmed mean, median):

- Treat large deviations as suspicious.
- Remove or downweight outliers.
- But minority updates can look like outliers even if honest.

# Fairness-Robustness Trade-offs

Joint Optimization Approaches Methods to achieve both robustness and fairness simultaneously.

## Fairness-aware FL may require:

- Personalized models
- Reweighting
- Clustering
- Multi-task learning
- Adaptive thresholds

## Proposed Solutions:

- Multi-objective optimization frameworks
- Fairness-aware adversarial training
- Regularization techniques for balanced performance

**Multi-objective optimization needed for both robustness and fairness!**

# Secure federated learning: Secure Aggregation

Secure aggregation ensures the server can compute the aggregate of client updates without learning anything about individual client contributions.

## Security Requirements:

1. Privacy: Server learns only the aggregate
2. Correctness: Honest inputs included correctly
3. Robustness: Works with malicious clients
4. Efficiency: Reasonable computation/communication

# Homomorphic Encryption:principle

Homomorphic encryption allows computation on encrypted data without decryption.

**Key Property (Additive HE):**

$$Enc(a) \cdot Enc(b) = Enc(a + b)$$

**Schemes:**

- Paillier: Probabilistic, additive homomorphic
- CKKS/BFV: Support floating-point operations

Google if you want to know more: homomorphuic encryption is an advanced topic in cryptography

# Homomorphic Encryption:principle

SecureAgg Protocol (Google) uses threshold homomorphic encryption in TensorFlow Federated.

## 4-Phase Protocol:

1. Setup: Clients generate keys, establish secure channels
2. Masking: Clients add random masks to updates
3. Aggregation: Server sums encrypted masked updates
4. Unmasking: Clients reveal masks to cancel randomness

# Homomorphic Encryption:Paillier Cryptosystem

- Key Gen:  $(n, g)$  where  $n = p \cdot q$
- Encryption:  $Enc(x) = g^x \cdot r^n \bmod n^2$
- Addition:  $Enc(x) \cdot Enc(y) = Enc(x + y)$

## FL Protocol:

1. Server generates Paillier key pair  $(pk, sk)$
2. Clients:  $c_k = Enc_{pk}(w_k)$
3. Server:  $C = \prod c_k = Enc_{pk}(\sum w_k)$
4. Server:  $\sum w_k = Dec_{sk}(C)$

